

# Apply SVM to Amazon reviews data set [M]

June 1, 2018

```
In [2]: !pip install --user imblearn
```

Collecting imblearn

Using cached <https://files.pythonhosted.org/packages/81/a7/4179e6ebfd654bd0eac0b9c06125b8b4c96>  
Collecting imbalanced-learn (from imblearn)

Using cached <https://files.pythonhosted.org/packages/80/a4/900463a3c0af082aed9c5a43f4ec317a946>  
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.6/dist-packages (from imbalanced-learn==0.3.3)  
Requirement already satisfied: scipy in /usr/local/lib/python3.6/dist-packages (from imbalanced-learn==0.3.3)  
Requirement already satisfied: numpy in /usr/local/lib/python3.6/dist-packages (from imbalanced-learn==0.3.3)  
Installing collected packages: imbalanced-learn, imblearn  
Successfully installed imbalanced-learn-0.3.3 imblearn-0.0

You are using pip version 9.0.1, however version 10.0.1 is available. You should consider upgrading using the [recommended command](#).

```
In [3]: from sklearn.model_selection import train_test_split
```

```
from sklearn.grid_search import GridSearchCV
from sklearn.grid_search import RandomizedSearchCV
from scipy.stats import uniform
from scipy.stats import norm
```

```
from imblearn.over_sampling import SMOTE
```

```
import sqlite3
```

```
import pandas as pd
import numpy as np
```

```
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.feature_extraction.text import CountVectorizer
import gensim
```

```
from sklearn.metrics import classification_report, accuracy_score, confusion_matrix
```

```
from sklearn.preprocessing import StandardScaler
```

```
/home/vishal/anaconda3/lib/python3.6/site-packages/sklearn/cross_validation.py:41: DeprecationWarning:
    "This module will be removed in 0.20.", DeprecationWarning)
```

```
/home/vishal/anaconda3/lib/python3.6/site-packages/sklearn/grid_search.py:42: DeprecationWarning
  DeprecationWarning)
```

```
In [17]: con = sqlite3.connect('final.sqlite') # this is cleaned dataset
        final = pd.read_sql_query("""
        SELECT Score, Text_not_included
        FROM reviews
        """, con)[:2000]
```

```
In [18]: for i, seq in enumerate(final['Text_not_included']):
        final['Text_not_included'][i]=final['Text_not_included'][i].decode('UTF-8')
        X_train, X_test, y_train , y_test = train_test_split(final['Text_not_included'], final[
```

## 0.1 Generate TF IDF weighted word2Vec

```
In [19]: tf_idf_vect=TfidfVectorizer(ngram_range=(1,2), min_df=10, dtype=float)
        tf_idf_vect.fit(X_train)
        tf_idf_train=tf_idf_vect.transform(X_train)
        tf_idf_test=tf_idf_vect.transform(X_test)
```

```
In [20]: w2vec_model=gensim.models.word2vec.Word2Vec(sentences, min_count=10)

        avg_w2vec_train=np.zeros(shape=(len(X_train), 100), dtype=float)
```

```
In [21]: sentences=[]
        for review in X_train:
            sentence=[]
            for word in review.split():
                sentence.append(word)
            sentences.append(sentence)

        tf_idf_w2vec_train=np.zeros((len(X_train), 100), dtype=float)
        feat=tf_idf_vect.get_feature_names()
        for i, sentence in enumerate(sentences):
            tf_idf_sum=0
            for word in sentence:
                try:
                    tf_idf_w2vec_train[i]+=w2vec_model.wv[word]*tf_idf_train[i, feat.index(word)]
                    tf_idf_sum+=tf_idf_train[i, feat.index(word)]
                except KeyError:
                    pass
                except ValueError:
                    pass
            tf_idf_w2vec_train[i]/=tf_idf_sum

        sentences=[]
        for review in X_test:
            sentence=[]
```

```

        for word in review.split():
            sentence.append(word)
        sentences.append(sentence)

tf_idf_w2vec_test=np.zeros((len(X_test), 100), dtype=float)

for i, sentence in enumerate(sentences):
    tf_idf_sum=0
    for word in sentence:
        try:
            tf_idf_w2vec_test[i]+=w2vec_model.wv[word]*tf_idf_test[i, feat.index(word)]
            tf_idf_sum+=tf_idf_test[i, feat.index(word)]
        except KeyError:
            pass
        except ValueError:
            pass
    tf_idf_w2vec_test[i]/=tf_idf_sum

In [22]: # Upsampling minority class
over_sampler = SMOTE(ratio='minority')

tf_idf_train_resampled, y_train_resampled = over_sampler.fit_sample(tf_idf_train, y_train)

tf_idf_w2vec_train_resampled, y_train_resampled = over_sampler.fit_sample(tf_idf_w2vec_train, y_train)

scaler_tf_idf=StandardScaler(with_mean=False)

scaler_tf_w2vec=StandardScaler()

scaler_tf_idf.fit(tf_idf_train_resampled)

scaler_tf_w2vec.fit(tf_idf_w2vec_train_resampled)

tf_idf_train_scaled=scaler_tf_idf.transform(tf_idf_train_resampled)

tf_idf_w2vec_train_scaled=scaler_tf_w2vec.transform(tf_idf_w2vec_train_resampled)

tf_idf_test_scaled=scaler_tf_idf.transform(tf_idf_test)

tf_idf_w2vec_test_scaled=scaler_tf_w2vec.transform(tf_idf_w2vec_test)

In [23]: from sklearn.svm import SVC

In [32]: tuned_parameters = {'C': np.linspace(10, 20, 20, dtype=float), 'gamma' : np.linspace(0.0, 1.0, 20, dtype=float)}

```

```

#Using GridSearchCV
gscv = GridSearchCV(SVC(), tuned_parameters, scoring = 'accuracy', cv=5)

print(gscv.fit(tf_idf_w2vec_train_scaled, y_train_resampled))

tuned_parameters = {'C' : uniform(10,20), 'gamma' : uniform(0, 1)}

#Using RandomizedSearchCV
rscv = RandomizedSearchCV(SVC(), tuned_parameters, scoring = 'accuracy', cv=5, n_iter=100)

print(rscv.fit(tf_idf_w2vec_train_scaled, y_train_resampled))

GridSearchCV(cv=5, error_score='raise',
             estimator=SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
                           decision_function_shape='ovr', degree=3, gamma='auto', kernel='rbf',
                           max_iter=-1, probability=False, random_state=None, shrinking=True,
                           tol=0.001, verbose=False),
             fit_params={}, iid=True, n_jobs=1,
             param_grid={'C': array([ 10.         , 10.52632, 11.05263, 11.57895, 12.10526, 12.63158,
                                     13.15789, 13.68421, 14.21053, 14.73684, 15.26316, 15.78947,
                                     16.31579, 16.84211, 17.36842, 17.89474, 18.42105, 18.94737,
                                     19.47368, 20.         ]), 'gamma': array([ 0.001   , 0.0535...937,
                                   0.63195, 0.68453, 0.73711, 0.78968, 0.84226, 0.89484,
                                   0.94742, 1.         ])}),
             pre_dispatch='2*n_jobs', refit=True, scoring='accuracy', verbose=0)

RandomizedSearchCV(cv=5, error_score='raise',
                   estimator=SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
                                 decision_function_shape='ovr', degree=3, gamma='auto', kernel='rbf',
                                 max_iter=-1, probability=False, random_state=None, shrinking=True,
                                 tol=0.001, verbose=False),
                   fit_params={}, iid=True, n_iter=10, n_jobs=1,
                   param_distributions={'C': <scipy.stats._distn_infrastructure.rv_frozen object at 0x7fe...>,
                                       'gamma': <scipy.stats._distn_infrastructure.rv_frozen object at 0x7fe...>},
                   pre_dispatch='2*n_jobs', random_state=None, refit=True,
                   scoring='accuracy', verbose=0)

In [33]: predictions = gscv.best_estimator_.predict(tf_idf_w2vec_test_scaled)
print(classification_report(y_test, predictions))
print(confusion_matrix(y_test, predictions).T)
tn, fp, fn, tp = confusion_matrix(y_test, predictions).ravel()

print("TPR = {}\n TNR = {}\n FPR = {}\n FNR = {}".format(tp/(fn+tp), tn/(tn+fp), fp/(tn+fp), fn/(fn+tp)))

predictions = rscv.best_estimator_.predict(tf_idf_w2vec_test_scaled)
print(classification_report(y_test, predictions))
print(confusion_matrix(y_test, predictions).T)
tn, fp, fn, tp = confusion_matrix(y_test, predictions).ravel()

print("TPR = {}\n TNR = {}\n FPR = {}\n FNR = {}".format(tp/(fn+tp), tn/(tn+fp), fp/(tn+fp), fn/(fn+tp)))

```

	precision	recall	f1-score	support
negative	0.57	0.05	0.09	87
positive	0.79	0.99	0.88	313
avg / total	0.74	0.79	0.71	400

```
[[ 4  3]
 [ 83 310]]
TPR = 0.9904153354632588
TNR = 0.04597701149425287
FPR = 0.9540229885057471
FNR = 0.009584664536741214
```

	precision	recall	f1-score	support
negative	0.21	0.03	0.06	87
positive	0.78	0.96	0.86	313
avg / total	0.66	0.76	0.69	400

```
[[ 3 11]
 [ 84 302]]
TPR = 0.9648562300319489
TNR = 0.034482758620689655
FPR = 0.9655172413793104
FNR = 0.03514376996805112
```

```
In [34]: gscv.best_estimator_
```

```
Out[34]: SVC(C=10.0, cache_size=200, class_weight=None, coef0=0.0,
  decision_function_shape='ovr', degree=3, gamma=0.63194736842105259,
  kernel='rbf', max_iter=-1, probability=False, random_state=None,
  shrinking=True, tol=0.001, verbose=False)
```

```
In [35]: rscv.best_estimator_
```

```
Out[35]: SVC(C=20.582199592274449, cache_size=200, class_weight=None, coef0=0.0,
  decision_function_shape='ovr', degree=3, gamma=0.49642755839081232,
  kernel='rbf', max_iter=-1, probability=False, random_state=None,
  shrinking=True, tol=0.001, verbose=False)
```

## 0.2 Conclusion

SVMs take a lot of time to train even on small datasets with less dimensions. Best gamma is about 0.5 TNR as found by GSCV as well as RSCV is very poor C can be taken to be between 10 and 20