

# Apply\_Naive\_Bayes\_to\_Amazon\_reviews\_[M]

May 27, 2018

```
In [0]: !pip install -U -q PyDrive
```

```
In [0]: from pydrive.auth import GoogleAuth
        from pydrive.drive import GoogleDrive
        from google.colab import auth
        from oauth2client.client import GoogleCredentials
```

```
# 1. Authenticate and create the PyDrive client.
auth.authenticate_user()
gauth = GoogleAuth()
gauth.credentials = GoogleCredentials.get_application_default()
drive = GoogleDrive(gauth)
```

```
In [3]: file_list = drive.ListFile({'q': "'1pbLvjcisi6UtFm3sPciCJGbCG4NK3uyuS' in parents and tra
        for file1 in file_list:
            print('title: %s, id: %s' % (file1['title'], file1['id']))
```

```
title: Apply Naive Bayes to Amazon reviews [M].ipynb, id: 1qPxAZeYQUM-eqaKnOSM5ubK2IPIVmdyo
title: clean_final.sqlite, id: 1TOHyUqaVFyD8HfIQEM6WN8jF8SpEOsAo
title: KNN on Credit Card fraud detection.ipynb, id: 1CkA-RBfXqvubKkQrpnjbYUKVsC7VH1T1
title: creditcard.csv, id: 1VpeqlS01PVrlz1MIqvQTzc3Pno_Cj4SV
title: creditcard.csv, id: 1bnZktEq3N_5wjoCH85oIXHxNwXUW_jx-
title: Untitled, id: 1K0wwkizWx3W08d-zw-YewWlUrPdINYmp
title: final.sqlite, id: 10zLc3k6-T55I-XRMq47ERyCbQbVw4caF
title: HeavyComputations.ipynb, id: 1aB0Re3gqeFY-iNhzMtr-TIkzEyEvFxcG
```

```
In [0]: sql = drive.CreateFile({'id': '10zLc3k6-T55I-XRMq47ERyCbQbVw4caF'})
        sql.GetContentFile('final.sqlite')
```

```
In [0]: %matplotlib inline
        import pandas as pd
        from sklearn.feature_extraction.text import TfidfVectorizer
        import numpy as np
        import matplotlib.pyplot as plt
        import sqlite3
```

```
In [0]: con = sqlite3.connect('final.sqlite') # this is cleaned dataset
        final = pd.read_sql_query("""
```

```

SELECT Score, Text_not_included
FROM reviews
""", con)

```

```
In [65]: len(final)
```

```
Out[65]: 23953
```

```
In [0]: for i, seq in enumerate(final['Text_not_included']):
        final['Text_not_included'][i]=final['Text_not_included'][i].decode('UTF-8')
```

```
In [0]: from sklearn.model_selection import train_test_split
```

```
In [0]: X_train, X_test, y_train , y_test = train_test_split(final['Text_not_included'], final['
```

```
In [0]: X_train, X_val, y_train, y_val = train_test_split(X_train, y_train, test_size=0.25, shuf
```

```
In [80]: type(y_train)
```

```
Out[80]: pandas.core.series.Series
```

```
In [12]: !pip install imblearn
```

```
Collecting imblearn
```

```
  Downloading https://files.pythonhosted.org/packages/81/a7/4179e6ebfd654bd0eac0b9c06125b8b4c96a
```

```
Collecting imbalanced-learn (from imblearn)
```

```
  Downloading https://files.pythonhosted.org/packages/80/a4/900463a3c0af082aed9c5a43f4ec317a9469
```

```
    100% || 153kB 7.2MB/s
```

```
Requirement already satisfied: numpy in /usr/local/lib/python3.6/dist-packages (from imbalanced-
```

```
Requirement already satisfied: scipy in /usr/local/lib/python3.6/dist-packages (from imbalanced-
```

```
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.6/dist-packages (from imba
```

```
Installing collected packages: imbalanced-learn, imblearn
```

```
Successfully installed imbalanced-learn-0.3.3 imblearn-0.0
```

```
In [0]: from sklearn.feature_extraction.text import CountVectorizer
```

```
In [100]: bbow = CountVectorizer(ngram_range = (1,2))
        bbow.fit(X_train)
```

```
Out[100]: CountVectorizer(analyzer='word', binary=False, decode_error='strict',
                        dtype=<class 'numpy.int64'>, encoding='utf-8', input='content',
                        lowercase=True, max_df=1.0, max_features=None, min_df=1,
                        ngram_range=(1, 2), preprocessor=None, stop_words=None,
                        strip_accents=None, token_pattern='(?u)\\b\\w\\w+\\b',
                        tokenizer=None, vocabulary=None)
```

```
In [0]: text_vectors = bbow.transform(X_train)
```

```

In [0]: # Oversampling train set
import imblearn
over_sampler=imblearn.over_sampling.SMOTE(ratio='minority')
X_train_resampled, y_train_resampled = over_sampler.fit_sample(text_vectors, y_train)

In [103]: X_train_resampled.get_shape()[0]==y_train_resampled.shape[0]

Out[103]: True

In [104]: len(y_train_resampled[np.where(y_train_resampled=='positive')])==len(y_train_resampled)

Out[104]: True

In [105]: X_train_resampled.getrow(5).todense()

Out[105]: matrix([[0., 0., 0., ..., 0., 0., 0.]])

In [0]: text_vectors_val=bbow.transform(X_val)


In [0]: from sklearn.naive_bayes import BernoulliNB
from sklearn.cross_validation import cross_val_score
from sklearn.metrics import classification_report, accuracy_score, confusion_matrix

In [114]: x=[]
y=[]
for alpha in range(1,5001,1):
    nb_clf = BernoulliNB(alpha=alpha, class_prior=[1,1])
    nb_clf.fit(X_train_resampled, y_train_resampled)
    predictions = nb_clf.predict(text_vectors_val)
    accuracy = accuracy_score(y_val, predictions)
    x.append(alpha)
    y.append(accuracy)

plt.plot(x, y, 'go')
print('max accuracy {} at alpha = {}'.format(np.max(y), x[np.argmax(y)]))

max accuracy 0.8714255896472552 at alpha = 2

```



```
In [115]: print('max accuracy {} at alpha = {}'.format(np.max(y), x[np.argmax(y)]))
```

```
max accuracy 0.8714255896472552 at alpha = 2
```

```
In [0]: test_vectors=bbow.transform(X_test)
```

```
In [0]: nb_clf = BernoulliNB(alpha=2, class_prior=[1,1])
nb_clf.fit(X_train_resampled, y_train_resampled)
predictions = nb_clf.predict(test_vectors)
accuracy = accuracy_score(y_test, predictions)
```

```
In [118]: accuracy
```

```
Out[118]: 0.8668336464203715
```

```
In [119]: print(classification_report(y_test, predictions))
```

	precision	recall	f1-score	support
negative	0.71	0.22	0.34	733
positive	0.88	0.98	0.93	4058
avg / total	0.85	0.87	0.84	4791

```

In [120]: print(confusion_matrix(y_test, predictions).T)

[[ 163   68]
 [ 570 3990]]

In [135]: nb_clf.classes_

Out[135]: array(['negative', 'positive'], dtype='<U8')

In [134]: X_train_resampled.shape

Out[134]: (23972, 303207)

In [138]: len(nb_clf.feature_count_[0])

Out[138]: 303207

In [0]: p_wi_given_y_negative = np.array(np.reshape(nb_clf.feature_log_prob_[0], [-1]))

In [0]: p_wi_given_y_positive = np.array(np.reshape(nb_clf.feature_log_prob_[1], [-1]))

In [0]: feat_indices_negative=np.argsort(-p_wi_given_y_negative)
        feat_indices_positive=np.argsort(-p_wi_given_y_positive)

In [154]: type(bbow.get_feature_names())

Out[154]: list

In [0]: important_features_negative = np.array(bbow.get_feature_names())[feat_indices_negative]
        important_features_positive = np.array(bbow.get_feature_names())[feat_indices_positive]

In [164]: important_features_negative[0:100]

Out[164]: array(['not', 'like', 'product', 'tast', 'would', 'buy', 'tri', 'get',
                  'one', 'price', 'good', 'flavor', 'dont', 'use', 'review', 'way',
                  'purchas', 'made', 'food', 'even', 'store', 'great', 'time',
                  'order', 'eat', 'dog', 'box', 'make', 'would not', 'much',
                  'bought', 'look', 'love', 'bag', 'think', 'amazon', 'first',
                  'realli', 'disappoint', 'could', 'better', 'also', 'didnt',
                  'foster', 'price foster', 'smith', 'foster smith', 'want', 'say',
                  'bad', 'littl', 'chip', 'okay', 'packag', 'know', 'give', 'well',
                  'expens', 'thought', 'got', 'never', 'differ', 'found', 'brand',
                  'thing', 'someth', 'not way', 'two', 'okay would', 'way buy',
                  'treat', 'ive', 'local', 'howev', 'away', 'quit', 'find', 'day',
                  'money', 'open', 'still', 'receiv', 'back', 'tast like', 'ingredi',
                  'recommend', 'year', 'sugar', 'product made', 'item', 'drink',
                  'old', 'mix', 'mani', 'great review', 'enjoy', 'contain',
                  'review product', 'per', 'lot'], dtype='<U33')

In [165]: important_features_positive[0:100]

```

```
Out[165]: array(['not', 'like', 'love', 'good', 'tast', 'great', 'flavor', 'one',
                'tri', 'use', 'product', 'get', 'make', 'food', 'buy', 'would',
                'time', 'eat', 'best', 'realli', 'amazon', 'also', 'much', 'find',
                'price', 'dont', 'well', 'littl', 'recommend', 'order', 'store',
                'dog', 'even', 'better', 'ive', 'year', 'bag', 'high', 'tea',
                'day', 'give', 'mix', 'treat', 'sweet', 'look', 'first', 'want',
                'found', 'favorit', 'chip', 'enjoy', 'drink', 'delici', 'think',
                'way', 'keep', 'work', 'thing', 'bit', 'brand', 'need', 'made',
                'lot', 'sinc', 'nice', 'know', 'sugar', 'purchas', 'pack',
                'bought', 'coffe', 'come', 'snack', 'say', 'could', 'perfect',
                'still', 'two', 'packag', 'add', 'everi', 'mani', 'alway',
                'without', 'never', 'healthi', 'differ', 'cant', 'take', 'seem',
                'local', 'box', 'easi', 'wonder', 'ever', 'water', 'right', 'got',
                'put', 'fresh'], dtype='<U33')
```

## 0.1 Results

Optimal alpha = 2 Top 100 features found most commonly in negative and positive reviews are enumerated above Accuracy on test set = 86.68% TPR (recall wrt positive class) = 98% TNR (recall wrt negative class) = 22% FPR = 78% FNR = 2% Precision wrt positive class = 0.88 Precision wrt negative class = 0.71 F1 score wrt positive class = 0.93 F1 score wrt negative class = 0.34

### 0.1.1 Conclusion

Even after resampling of minority class using SMOTE we observe that Naive Bayes classifier fails miserably in distinguishing negative class from positive class. The ill effects of majority class (positive reviews) dominance are clearly visible in TNR above. More than 50% of negative reviews are being classified as positive. This clearly shows that the 'naive' approximation although giving satisfactory results in spam filters, is failing in this case and hence this classifier is not appropriate for text classifications in general.